

Quickstart

February 19, 2025

1 Quickstart Guide to Pytorch

This section runs through the API for common tasks in machine learning. Generally this section is for someone who is familiar with other deep learning frameworks.

1.1 Working with Data

PyTorch has two primitives to work with data: `torch.utils.data.DataLoader` and `torch.utils.data.Dataset`. `Dataset` stores the samples and their corresponding labels, and `DataLoader` wraps an iterable around the `Dataset`.

```
[1]: import torch
      from torch import nn
      from torch.utils.data import DataLoader
      from torchvision import datasets
      from torchvision.transforms import ToTensor
```

PyTorch offers domain-specific libraries such as `TorchText`, `TorchVision`, and `TorchAudio`, all of which include datasets. For this tutorial, we will be using a `TorchVision` dataset.

The `torchvision.datasets` module contains `Dataset` objects for many real-world vision data like CIFAR, COCO. In this tutorial, we use the FashionMNIST dataset. Every `TorchVision Dataset` includes two arguments: `transform` and `target_transform` to modify the samples and labels respectively.

```
[2]: # Download training data from open datasets.
      training_data = datasets.FashionMNIST(
          root="data",
          train=True,
          download=True,
          transform=ToTensor(),
      )

      # Download test data from open datasets.
      test_data = datasets.FashionMNIST(
          root="data",
          train=False,
          download=True,
          transform=ToTensor(),
```

```
)
```

We pass the `Dataset` as an argument to `DataLoader`. This wraps an iterable over our dataset, and supports automatic batching, sampling, shuffling and multiprocess data loading. Here we define a batch size of 64, i.e. each element in the dataloader iterable will return a batch of 64 features and labels.

```
[3]: batch_size = 64

# Create data loaders.
train_dataloader = DataLoader(training_data, batch_size=batch_size)
test_dataloader = DataLoader(test_data, batch_size=batch_size)

for X, y in test_dataloader:
    print(f"Shape of X [N, C, H, W]: {X.shape}")
    print(f"Shape of y: {y.shape} {y.dtype}")

    # Time to introspect further
    print(f"dtype of X: {X.dtype}") # Tells us the type of data in the input
    ↪data
    print(f"The min and max values in X respectively: {X.min()}, {X.max()}")
    break
```

```
Shape of X [N, C, H, W]: torch.Size([64, 1, 28, 28])
```

```
Shape of y: torch.Size([64]) torch.int64
```

```
dtype of X: torch.float32
```

```
The min and max values in X respectively: 0.0, 1.0
```

Furthermore I felt like extracting and looking at a single image so here is what I did

```
[4]: import matplotlib.pyplot as plt
# import numpy as np

single_image, single_label = training_data[6]
single_image = single_image.squeeze().numpy() # Remove the batch dimension and
    ↪convert to NumPy
# single_image = (single_image * 255).astype(np.uint8) # Scale to [0, 255] and
    ↪convert to uint8

plt.imshow(single_image, cmap='gray')
plt.title(f"Label: {single_label}")
plt.axis('off') # Hide the axes
plt.show()
# plt.imsave('example.jpg', single_image, cmap='gray')
```

Label: 7



We can also inspect multiple images, for example create a grid of a batch in the input data

```
[5]: import torchvision.utils as vutils

# Create a grid of images
grid = vutils.make_grid(X, nrow=8, normalize=True, scale_each=True)

# Convert the grid to a NumPy array and permute dimensions
grid = grid.permute(1, 2, 0).numpy()

# Display the grid
plt.imshow(grid)
plt.axis('off')
plt.show()
```



1.2 Creating Models

To define a neural network in PyTorch, we create a class that inherits from `nn.Module`. We define the layers of the network in the `__init__` function and specify how data will pass through the network in the `forward` function. To accelerate operations in the neural network, we move it to the `accelerator` such as CUDA, MPS, MTIA, or XPU. If the current accelerator is available, we will use it. Otherwise, we use the CPU.

EDIT: The above information regarding accelerators seems outdated as `torch.accelerator` only supports `CUDA`, `MTIA`, `XPU`.

For mac systems we will instead directly check `MPS` availability or use `cpu` instead.

```
[6]: # For CUDA, MTIA, or XPU
# device = torch.accelerator.current_accelerator().type if torch.accelerator.
# is_available() else torch.device("cpu")
# print(f"Using {device} device")

# For MPS (Mac systems)
device = torch.device("mps") if torch.backends.mps.is_available() else torch.
device("cpu")
print(f"Using {device} device")
```

Using mps device

```
[7]: # Define model
class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10)
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

model = NeuralNetwork().to(device)
print(model)
print(f"Model is on device: {next(model.parameters()).device}") # Should print_
↪ 'mps'
```

```
NeuralNetwork(
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (linear_relu_stack): Sequential(
    (0): Linear(in_features=784, out_features=512, bias=True)
    (1): ReLU()
    (2): Linear(in_features=512, out_features=512, bias=True)
    (3): ReLU()
    (4): Linear(in_features=512, out_features=10, bias=True)
  )
)
Model is on device: mps:0
```

1.3 Optimising the Model Parameters

To train a model, we need a [loss function](#) and an [optimiser](#).

```
[8]: loss_fn = nn.CrossEntropyLoss().to(device)
optimizer = torch.optim.SGD(model.parameters(), lr=1e-3)
```

In a single training loop, the model makes predictions on the training dataset (fed to it in batches), and backpropagates the prediction error to adjust the model's parameters.

```
[9]: def train(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    model.train()
```

```

for batch, (X, y) in enumerate(dataloader):
    X, y = X.to(device), y.to(device)

    # Compute prediction error
    pred = model(X)
    loss = loss_fn(pred, y)

    # Backpropagation
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()

    if batch % 100 == 0:
        loss, current = loss.item(), (batch + 1) * len(X)
        print(f"loss: {loss:>7f} [{current:>5d}/{size:>5d}]")

```

We also check the model's performance against the test dataset to ensure it is learning.

```

[10]: def test(dataloader, model, loss_fn):
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    model.eval()
    test_loss, correct = 0, 0
    with torch.no_grad():
        for X, y in dataloader:
            X, y = X.to(device), y.to(device)
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()
    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss:␣
↪{test_loss:>8f} \n")

```

The training process is conducted over several iterations (epochs). During each epoch, the model learns parameters to make better predictions. We print the model's accuracy and loss at each epoch; we'd like to see the accuracy increase and the loss decrease with every epoch.

```

[11]: epochs = 5
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train(train_dataloader, model, loss_fn, optimizer)
    test(test_dataloader, model, loss_fn)
print("Done!")

```

Epoch 1

```

-----
loss: 2.311492 [ 64/60000]
loss: 2.293662 [ 6464/60000]

```

loss: 2.278553 [12864/60000]
loss: 2.269704 [19264/60000]
loss: 2.252427 [25664/60000]
loss: 2.227463 [32064/60000]
loss: 2.240565 [38464/60000]
loss: 2.201384 [44864/60000]
loss: 2.198047 [51264/60000]
loss: 2.173079 [57664/60000]

Test Error:

Accuracy: 50.9%, Avg loss: 2.166078

Epoch 2

loss: 2.174698 [64/60000]
loss: 2.158323 [6464/60000]
loss: 2.109015 [12864/60000]
loss: 2.128747 [19264/60000]
loss: 2.073018 [25664/60000]
loss: 2.021720 [32064/60000]
loss: 2.061239 [38464/60000]
loss: 1.972798 [44864/60000]
loss: 1.973918 [51264/60000]
loss: 1.918002 [57664/60000]

Test Error:

Accuracy: 54.5%, Avg loss: 1.907859

Epoch 3

loss: 1.933078 [64/60000]
loss: 1.898628 [6464/60000]
loss: 1.790941 [12864/60000]
loss: 1.838273 [19264/60000]
loss: 1.721255 [25664/60000]
loss: 1.676981 [32064/60000]
loss: 1.718241 [38464/60000]
loss: 1.602537 [44864/60000]
loss: 1.621176 [51264/60000]
loss: 1.528310 [57664/60000]

Test Error:

Accuracy: 61.1%, Avg loss: 1.536844

Epoch 4

loss: 1.598383 [64/60000]
loss: 1.557854 [6464/60000]
loss: 1.412430 [12864/60000]
loss: 1.488797 [19264/60000]
loss: 1.359946 [25664/60000]

```
loss: 1.362863 [32064/60000]
loss: 1.395155 [38464/60000]
loss: 1.304919 [44864/60000]
loss: 1.333764 [51264/60000]
loss: 1.241617 [57664/60000]
Test Error:
  Accuracy: 63.5%, Avg loss: 1.262457
```

Epoch 5

```
-----
loss: 1.338601 [  64/60000]
loss: 1.313664 [ 6464/60000]
loss: 1.152314 [12864/60000]
loss: 1.257326 [19264/60000]
loss: 1.126186 [25664/60000]
loss: 1.161881 [32064/60000]
loss: 1.198766 [38464/60000]
loss: 1.122901 [44864/60000]
loss: 1.157875 [51264/60000]
loss: 1.077889 [57664/60000]
Test Error:
  Accuracy: 64.8%, Avg loss: 1.094102
```

Done!

More info on [training your model](#).

1.4 Saving Models

A common way to save a model is to serialize the internal state dictionary (containing the model parameters).

```
[12]: torch.save(model.state_dict(), "model.pth")
      print("Saved PyTorch Model State to model.pth")
```

Saved PyTorch Model State to model.pth

1.5 Loading Models

The process for loading a model includes re-creating the model structure and loading the state dictionary into it.

```
[13]: model = NeuralNetwork().to(device)
      model.load_state_dict(torch.load("model.pth", weights_only=True))
```

```
[13]: <All keys matched successfully>
```

This model can now be used to make predictions.


```
[14]: classes = [
        "T-shirt/top",
        "Trouser",
        "Pullover",
        "Dress",
        "Coat",
        "Sandal",
        "Shirt",
        "Sneaker",
        "Bag",
        "Ankle boot",
    ]

model.eval()
x, y = test_data[0][0], test_data[0][1]
with torch.no_grad():
    x = x.to(device)
    pred = model(x)
    predicted, actual = classes[pred[0].argmax(0)], classes[y]
    print(f'Predicted: "{predicted}", Actual: "{actual}"')
```

Predicted: "Ankle boot", Actual: "Ankle boot"

More info on [saving and loading your model](#).